

EMPLOYMENT	Allen Institute for AI <i>Research Scientist</i> • Advisor: Yejin Choi	Seattle, US 2023 - 2024
	Allen Institute for AI <i>Postdoctoral Fellow</i> • Advisor: Yejin Choi	Seattle, US 2022 - 2023
	Mila – Quebec Artificial Intelligence Institute / McGill University <i>Visiting Scholar</i> • Advisor: Siva Reddy	Montreal, CA 2021 - 2022
	Google DeepMind <i>Student Researcher</i> • Advisors: Tal Linzen, David Reitter, Hannah Rashkin	NYC, US 2020 - 2022
	Microsoft Research <i>Research Intern</i> • Advisors: Alessandro Sordoni, Goeff Gordon	Montreal, CA 2019 - 2020
	Google DeepMind <i>Research Intern</i> • Advisors: Diyi Yang, Tom Kwiatkowski	NYC, US 2019
EDUCATION	Ph.D. Computing Science, University of Alberta <i>Thesis: Mitigating Hallucinations in Conversational LLMs.</i> • Advisor: Osmar Zaiane, GPA: 4.00/4.00	Edmonton, Canada 2018 - 2022
	MSc. Computing Science, University of Alberta <i>Thesis: Response Generation For An Open-Ended Conversational Agent</i> • Advisor: Osmar Zaiane, GPA: 4.00/4.00	Edmonton, Canada 2016 - 2018
PREPRINTS	You can find an exhaustive list of my publications in my Google Scholar profile.	
	1. RL Grokking Recipe: How Does RL Unlock and Transfer New Algorithms in LLMs? . Yiyu Sun, Yuhao Cao, Pohao Huang, Haoyue Bai, Hannaneh Hajishirzi, <u>Nouha Dziri</u> *, Dawn Song* (* = equal advising role) <i>Arxiv 2025 (in submission to ICLR)</i>	
	2. Climbing the Ladder of Reasoning: What LLMs Can-and Still Can't-Solve after SFT? . Yiyu Sun, Shawn Hu, Georgia Zhou, Ken Zheng, Hannaneh Hajishirzi, <u>Nouha Dziri</u> , Dawn Song. <i>Arxiv 2025</i>	
	3. OpenAgentSafety: A Comprehensive Framework for Evaluating Real-World AI Agent Safety. Sanidhya Vijayvargiya, Aditya Bharat Soni, Xuhui Zhou, Zora Zhiruo Wang, <u>Nouha Dziri</u> , Graham Neubig, Maarten Sap. <i>Arxiv 2025 (in submission to ICLR)</i>	

4. **OMEGA: Can LLMs Reason Outside the Box in Math? Evaluating Exploratory, Compositional, and Transformative Generalization.** Yiyu Sun, Georgia Zhou, Hao Wang, Dacheng Li, Nouha Dziri^{*}, Dawn Song^{*} (* = equal advising role)
NeurIPS 2025
5. **Artificial Hivemind: The Open-Ended Homogeneity of Language Models (and Beyond)** Liwei Jiang, Yuanjun Chai, Margaret Li, Mickel Liu, Raymond Fok, Maarten Sap, Yulia Tsvetkov, Nouha Dziri, Yejin Choi.
NeurIPS 2025
Oral presentation (Top 1.4% among 25k submissions)
6. **Why and How LLMs Hallucinate: Connecting the Dots with Subsequence Associations** Yiyu Sun, Yu Gai, Lijie Chen, Abhi Ravichander, Yejin Choi, Nouha Dziri, Dawn Song.
NeurIPS 2025
7. **2 OLMo 2 Furious.** Pete Walsh, Luca Soldaini, Dirk Groeneveld, Kyle Lo, Shane Arora, ..., Nouha Dziri, Noah A. Smith, Hannaneh Hajishirzi.
COLM 2025
2M downloads on HuggingFace
8. **TÜLU 3: Pushing Frontiers in Open Language Model Post-Training.** Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze Brahman, Lester James V. Miranda, Alisa Liu, Nouha Dziri, Shane Lyu, Yuling Gu, Saumya Malik, Victoria Graf, Jena D. Hwang, Jiangjiang Yang, Ronan Le Bras, Oyvind Tafjord, Chris Wilhelm, Luca Soldaini, Noah A. Smith, Yizhong Wang, Pradeep Dasigi and Hannaneh Hajishirzi.
COLM 2025
300K downloads on HuggingFace
9. **SafetyAnalyst: Interpretable, Transparent, and Steerable LLM Safety Moderation** Jing-Jing Li, Valentina Pyatkin, Max Kleiman-Weiner, Liwei Jiang, Nouha Dziri, Anne GE Collins, Jana Schaich Borg, Maarten Sap, Yejin Choi, Sydney Levine.
ICML 2025
10. **AI as Humanity’s Salieri: Quantifying Linguistic Creativity of Language Models via Systematic Attribution of Machine Text against Web Text.** Ximing Lu, Melanie Sclar, Skyler Hallinan, Niloofar Mireshghallah, Jiacheng Liu, Seungju Han, Allyson Ettinger, Liwei Jiang, Khyathi Chandu, Nouha Dziri, Yejin Choi.
ICLR 2025
Oral presentation (Top 1.8% among 12k submissions)
Featured in Science Magazine
11. **WildBench: Benchmarking LLMs with Challenging Tasks from Real Users in the Wild.** Bill Yuchen Lin, Yuntian Deng, Khyathi Chandu, Faeze Brahman, Abhilasha Ravichander, Valentina Pyatkin, Nouha Dziri, Ronan Le Bras, Yejin Choi.
ICLR 2025
12. **Steering Masked Discrete Diffusion Models Via Discrete Denoising Posterior Prediction** Jarrod Rector-Brooks, Mohsin Hasan, Zhangzhi Peng, Zachary Quinn, Chenghao Liu, Sarthak Mittal, Nouha Dziri, Michael Bronstein, Yoshua Bengio, Pranam Chatterjee, Alexander Tong, Avishek Joey Bose.
ICLR 2025
13. **RewardBench: Evaluating Reward Models for Language Modeling.** Nathan Lambert, Valentina Pyatkin, Jacob Morrison, LJ Miranda, Bill Yuchen Lin, Khyathi Chandu, Nouha Dziri, Sachin Kumar, Tom Zick, Yejin Choi, Noah A Smith, Hannaneh Hajishirzi.
NAACL 2025
14. **Rel-AI: An Interaction-Centered Approach To Measuring Human-LM Reliance.** Kaitlyn Zhou, Jena D Hwang, Xiang Ren, Nouha Dziri, Dan Jurafsky, Maarten Sap.
NAACL 2025
Best Paper Runner-Up 2025 (Top 0.2%)
15. **To Err is AI: A Case Study Informing LLM Flaw Reporting Practices** Sean McGregor, Allyson Ettinger, Nick Judd, Paul Albee, Liwei Jiang, Kavel Rao, Will Smith, Shayne Longpre, Avijit Ghosh, Christopher Fiorelli, Michelle Hoang, Sven Cattell, Nouha Dziri.
AAAI 2025

16. **WildTeaming at Scale: From In-the-Wild Jailbreaks to (Adversarially) Safer Language Models.** Liwei Jiang, Kavel Rao, Seungju Han, Allyson Ettinger, Faeze Brahman, Sachin Kumar, Niloofar Mireshghallah, Maarten Sap, Yejin Choi, Nouha Dziri.
NeurIPS 2024
17. **WildGuard: Open One-Stop Moderation Tools For Safety Risks, Jailbreaks, and Refusals of LLMs.** Seungju Han, Kavel Rao, Allyson Ettinger, Liwei Jiang, Bill Yuchen Lin, Nathan Lambert, Yejin Choi, and Nouha Dziri.
NeurIPS 2024
215K downloads on HuggingFace
18. **Multi-Attribute Constraint Satisfaction via Language Model Rewriting.** Ashutosh Baheti, Debanjana Chakraborty, Faeze Brahman, Ronan Le Bras, Ximing Lu, Nouha Dziri, Yejin Choi, Mark Riedl, Maarten Sap.
TMLR 2024
19. **A Roadmap to Pluralistic Alignment.** Taylor Sorensen, Jared Moore, Jillian Fisher, Mitchell Gordon, Niloofar Mireshghallah, Christopher Michael Rytting, Andre Ye, Liwei Jiang, Ximing Lu, Nouha Dziri, Tim Althoff, Yejin Choi.
ICML 2024
20. **Elastic Weight Removal For Faithful and Abstractive Dialogue generation.** Nico Daheim, Nouha Dziri, Mrinmaya Sachan, Iryna Gurevych, Edoardo Ponti.
NAACL 2024
21. **Faith and Fate: Limits of Transformers on Compositionality.** Nouha Dziri, Ximing Lu, Melanie Sclar, Xiang Lorraine Li, Liwei Jian, Bill Yuchen Lin, Peter West, Chandra Bhagavatula, Ronan Le Bras, Jena D Hwang, Soumya Sanyal, Sean Welleck, Xiang Ren, Allyson Ettinger, Zaid Harchaoui, Yejin Choi.
NeurIPS 2023
Spotlight (Top 2% among 12k submissions)
Featured in Quanta Magazine, Science News, Montreal AI Ethics Institute
15 invited talks, 1 podcast, 2 guest lectures
22. **Fine-Grained Human Feedback Gives Better Rewards for Language Model Training.** Zeqiu Wu, Yushi Hu, Weijia Shi, Nouha Dziri, Alane Suhr, Prithviraj Ammanabrolu, Noah A Smith, Mari Ostendorf, Hannaneh Hajishirzi.
NeurIPS 2023
Spotlight (Top 2% among 12k submissions)
23. **Self-Refine: Iterative Refinement with Self-Feedback.** Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, Sean Welleck, Bodhisattwa Prasad Majumder, Shashank Gupta, Amir Yazdanbakhsh, Peter Clark.
NeurIPS 2023
24. **The Generative AI Paradox:” What It Can Create, It May Not Understand”. Peter West*, Ximing Lu*, Nouha Dziri*, Faeze Brahman*, Linjie Li*, Jena D Hwang, Liwei Jiang, Jillian Fisher, Abhilasha Ravichander, Khyathi Chandu, Benjamin Newman, Pang Wei Koh, Allyson Ettinger, Yejin Choi. (* = equal contribution)**
ICLR 2024
25. **Phenomenal Yet Puzzling: Testing Inductive Reasoning Capabilities of Language Models with Hypothesis Refinement.** Linlu Qiu, Liwei Jiang, Ximing Lu, Melanie Sclar, Valentina Pyatkin, Chandra Bhagavatula, Bailin Wang, Yoon Kim, Yejin Choi, Nouha Dziri*, Xiang Ren*. (* = equal advising role)
ICLR 2024
Oral presentation (Top 1.8% among 7k submissions)
26. **The Unlocking Spell on Base LLMs: Rethinking Alignment via In-Context Learning.** Bill Yuchen Lin, Abhilasha Ravichander, Ximing Lu, Nouha Dziri, Melanie Sclar, Khyathi Chandu, Chandra Bhagavatula, Yejin Choi.
ICLR 2024

27. **Value Kaleidoscope: Engaging AI with pluralistic human values, rights, and duties**
Taylor Sorensen, Liwei Jiang, Jena Hwang, Sydney Levine, Valentina Pyatkin, Peter West, Nouha Dziri, Ximing Lu, Kavel Rao, Chandra Bhagavatula, Maarten Sap, John Tasioulas, Yejin Choi.
AAAI 2024
28. **Culture-Gen: Revealing global cultural perception in language models through natural language prompting**
Huihan Li, Liwei Jiang, Jena D Hwang, Hyunwoo Kim, Sebastin Santy, Taylor Sorensen, Bill Yuchen Lin, Nouha Dziri, Xiang Ren, Yejin Choi.
COLM 2024
29. **What Makes it Ok to Set a Fire? Iterative Self-distillation of Contexts and Rationales for Disambiguating Defeasible Social and Moral Situations** Kavel Rao, Liwei Jiang, Valentina Pyatkin, Yuling Gu, Niket Tandon, Nouha Dziri, Faeze Brahman, Yejin Choi.
EMNLP 2023
30. **Evaluating Open-Domain Question Answering in the Era of Large Language Models.**
Ehsan Kamalloo, Nouha Dziri, Charles LA Clarke, Davood Rafiei.
ACL 2023
31. **CHAMPAGNE: Learning Real-world Conversation from Large-Scale Web Videos.**
Seungju Han, Jack Hessel, Nouha Dziri, Yejin Choi, Youngjae Yu.
ICCV 2023
32. **On the Origin of Hallucinations in Conversational Models: Is it the Datasets or the Models?** Nouha Dziri, Sivan Milton, Mo Yu, Osmar Zaiane, Siva Reddy.
NAACL 2022
33. **FaithDial: A Faithful Benchmark for Information-Seeking Dialogue.**
Nouha Dziri, Ehsan Kamalloo, Sivan Milton, Osmar Zaiane, Mo Yu, Edoardo M Ponti, Siva Reddy.
TACL 2022
34. **Neural Path Hunter: Reducing Hallucination in Dialogue Systems via Path Grounding.**
Nouha Dziri, Andrea Madotto, Osmar Zaiane, Avishek Joey Bose.
EMNLP 2021
35. **Decomposed Mutual Information Estimation For Contrastive Representation Learning.**
Alessandro Sordani*, Nouha Dziri*, Hannes Schulz*, Geoff Gordon, Philip Bachman, Remi Tachet Des Combes. (* = equal contribution)
ICML 2021
36. **Evaluating Coherence in Dialogue Systems Using entailment.**
Nouha Dziri*, Ehsan Kamalloo*, Kory W Mathewson, Osmar Zaiane. (* = equal contribution)
NAACL 2019

Computer Use Agents

ICML workshop, Vancouver

July 2025

Panelists: Graham Neubig, Ruslan Salakhutdinov, Yu Su, Victor Zhong, Alexandre Drouin

Data in Generative Models

ICML workshop, Vancouver

July 2025

Panelists: Tatsunori Hashimoto, Eric Wong, Serena Booth, Pin-Yu Chen, Ivan Evtimov

AI Reasoning

Cross Future AI & Technology Summit, Vancouver

July 2025

Panelists: Wenhu Chen, Ken Perlin

Trustworthy Foundation Models

International Symposium on Trustworthy Foundation Models, UAE

May 2025

Panelists: Tomas Mikolov, Hakim Hacid, Tongliang Liu

INVITED PANELS	Making the Right Career Call across Academia and Industry <i>ICLR Workshop Women in ML, Singapore</i> Panelists: Claire Vernade, Reyhane Askari, Katherine Driscoll	April 2025
	Meta-Generation Algorithms for Large Language Models <i>NeurIPS Tutorial, Vancouver</i> Panelists: Noam Brown, Jakob Foerster, Rishabh Agarwal, Beidi Chen	Dec 2024
	Analyzing the Plasticity of LLMs During Language Interactions. <i>NeurIPS Language Gamification Workshop, Vancouver</i> Panelists: Aaron Courville, Alane Suhr, Tom Schaul, Marc Lanctot, Tom Griffiths, and Sam Devlin	Dec 2024
	Towards A Global Standard for AI Product Safety <i>MLCommons, San Francisco</i> Panelists: Ion Stoica, April Chen, Wan Sie Lee	Nov 2024
INVITED TALKS	Reasoning in LLMs <i>Workshop on Cognitive Basis of Reasoning/IVADO, Montreal</i>	Jan 2026
	Autonomous LLM Agents: Risks and Scientific Challenges <i>IVADO/Mila - Quebec Artificial Intelligence Institute, Montreal</i>	Nov 2025
	Faith and Fate: Limits of Transformers in Reasoning <i>D.E.Shaw Research, NYC</i>	Sept 2025
	LLM Reasoning: Advanced Inference-time Strategies <i>Armenia LLM Summer School, Armenia</i>	July 2025
	Can LLMs Reason Outside the Box in Math? <i>Apple Workshop on Reasoning and Planning, Cupertino</i>	July 2025
	Faith and Fate: Limits of Transformers in Reasoning <i>Cross Future AI & Technology Summit, Vancouver</i>	July 2025
	Trustworthy LLMs: How Data Quality Shapes Performance and Where It Falls Short? <i>ICML 2025 Workshop on Data in Generative Models, Vancouver</i>	July 2025
	OpenAgentSafety: A Comprehensive Framework for Evaluating AI Agent Safety <i>ICML 2025 Workshop on Computer-Use Agents, Vancouver</i>	July 2025
	Does Scaling Guarantee Trustworthy LLMs <i>International Symposium on Trustworthy Foundation Models, UAE</i>	May 2025
	In-Context Learning in LLMs: Potential and Limits <i>Causality in the Era of Foundation Models Workshop, Barbados</i>	Feb 2025
	Red-teaming and Safeguarding LLMs <i>International Association for Safe and Ethical AI, Paris</i>	Feb 2025
	In-Context Learning in LLMs: Potential and Limits <i>NeurIPS Language Gamification Workshop, Vancouver</i>	Dec 2024
	What it can create, it may not understand: Studying the Limits of Transformers <i>University of Cambridge</i>	May 2024
	Limits of Generative AI Models and their Societal Implications. <i>Princeton University</i>	Dec 2023
	Faith and Fate: Limits of Transformers on Compositionality <i>The Alan Turing Institute, UK</i>	Nov 2023

INVITED TALKS	Faith and Fate: Limits of Transformers on Compositionality. <i>University of Edinburgh</i>	Nov 2023
	Faith and Fate: Limits of Transformers on Compositionality <i>SAIL workshop on fundamental limits of LLMs, Germany</i>	Oct 2023
	Faith and Fate: Limits of Transformers on Compositionality <i>University of Pittsburgh</i>	Oct 2023
	Faith and Fate: Limits of Transformers on Compositionality <i>Formal Languages and Neural Networks Seminar, US</i>	Sep 2023
	Towards Building Hallucination-Free Conversational Models <i>Stanford University</i>	Aug 2022
	FaithDial: A Faithful Benchmark for Information-Seeking Dialogue <i>Google Research, NYC</i>	May 2022
	FaithDial: A Faithful Benchmark for Information-Seeking Dialogue <i>Amazon Research, Seattle</i>	June 2022
	Evaluating Coherence in Dialogue Systems Using Entailment. <i>Google DeepMind, Montreal</i>	Dec 2019
GUEST LECTURER	11-430/830 Ethics, Safety, and Social Impact in NLP and LLMs <i>Instructor: Maarten Sap</i> <i>Carnegie Mellon University</i> <i>Lecture: Redteaming and Safeguarding in LLMs</i>	Winter 2025
	Generative AI Seminar Course <i>Instructor: Adji Bousso Dieng</i> <i>Princeton University</i> <i>Lecture: Limits of Generative AI Models and their Societal Implications.</i>	Fall 2023
	CSE 599 D1: (Grad) Exploration on Language, Knowledge, and Reasoning <i>Instructor: Yejin Choi</i> <i>University of Washington</i> <i>Lecture: Can LLMs reason?</i>	Winter 2023
	CMPUT 650: (Grad) Computational Semantics <i>Instructor: Greg Kondrak</i> <i>University of Alberta</i> <i>Lecture: Challenges in Conversational LLMs</i>	Winter 2018
STUDENTS MENTORING	• Ximing Lu , Predoctoral student at Ai2 → PhD student at UW <i>ICLR 2025 Oral (top 1%), NeurIPS 2023 Spotlight (top 2%)</i>	2022-2024
	• Liwei Jiang , PhD student at UW <i>NeurIPS 2025 Oral (top 1%), NeurIPS 2024</i>	2023-2025
	• Seungju Han , BSc at Seoul National University → PhD at Stanford <i>NeurIPS 2024, ICCV 2023</i>	2023-2024
	• Kavel Rao , MSc student at UW → Researcher at Jane Street <i>NeurIPS 2024, EMNLP 2023</i>	2023-2024
	• Linlu Qiu , PhD student at MIT <i>ICLR 2024 Oral (top 1%)</i>	2023
	• Huihan Li , PhD student at University of Southern California <i>COLM 2024</i>	2023-2024

STUDENTS MENTORING	<ul style="list-style-type: none"> • Yiyu Sun, Postdoc at University of California, Berkeley <i>2 papers at NeurIPS 2025</i> 	2025-current
	<ul style="list-style-type: none"> • Shrey Jain, MSc student at Carnegie Mellon University 	2025-current
WORKSHOP Co- ORGANIZER	<ul style="list-style-type: none"> • The first workshop for Research on Agent Language Models (REALM) 	ACL 2025
	<ul style="list-style-type: none"> • The Reasoning and Planning for Large Language Models • System 2 Reasoning at Scale • The third workshop on Document-Grounded Dialogue Systems (DialDoc) 	ICLR 2025 NeurIPS 2024 ACL 2023
ACADEMIC SERVICES	Demo Chair: NAACL 2025	
	Senior Area Chair: ACL 2025 in the area of Ethics, Bias, and Fairness Area Chair: EMNLP 2023, COLM (2024-2025) Reviewer: NeurIPS (2022-2024), ICLR (2022-2024), ACL (2018-2023), EMNLP (2018-2023), NAACL (2018-2023), EACL (2018-2022)	
PRESS	<ul style="list-style-type: none"> • Quanta Magazine: <i>Chatbot Software Begins to Face Fundamental Limitations</i>, 2025. • Science Magazine: <i>AI writing is improving, but it can't match human creativity</i>, 2024. • LeMonde: <i>Should we be concerned about the 'hallucinations' of AI like ChatGPT or Gemini?</i>, 2024. • ScienceNews: <i>AI's reasoning skills can't be assessed by current tests</i>, 2024. • TechCrunch: <i>Treating a chatbot nicely might boost its performance — here's why</i>, 2024. • Podcast (Women in AI Research): <i>Limits of Transformers</i>, 2025. 	
AWARDS AND HONORS	<ul style="list-style-type: none"> • Runner-up Best Paper at NAACL 2025 (top 0.1%) 	2025
	<ul style="list-style-type: none"> • Outstanding Reviewer at ACL 2021 (top 1%) • Alberta Doctoral Recruitment Scholarship (\$10,000) • Mitacs Globalink PhD Graduate Fellowship (\$44,000) • National Scholarship for MSc and PhD studies (\$70,000) • Best Poster Award, ACM Canadian Celebration of Women in Computing (\$400) • Mitacs Globalink MSc Graduate Fellowship (\$15,000) • DAAD Scholarship for Research Internship, Leipzig, Germany (€10.000) • Erasmus Mundus Exchange Scholarship for BSc studies(€50.000) 	2021 2018 2018 2016 2017 2016 2015 2015