Nouha Dziri Ph.D.

nouha.dziri@gmail.com
http://nouhadziri.com
+1 206 617 2801

Employment

Allen Institute for AI   Seattle, US
Research Scientist   2023 - 2024
- Advisor: Yejin Choi

Allen Institute for AI   Seattle, US
Postdoctoral Fellow   2022 - 2023
- Advisor: Yejin Choi

Mila – Quebec Artificial Intelligence Institute / McGill University   Montreal, CA
Visiting Scholar   2021 - 2022
- Advisor: Siva Reddy

Google DeepMind   NYC, US
Student Researcher   2020 - 2022
- Advisors: Tal Linzen, David Reitter, Hannah Rashkin

Microsoft Research   Montreal, CA
Research Intern   2019 - 2020
- Advisors: Alessandro Sordoni, Goeff Gordon

Google DeepMind   NYC, US
Research Intern   2019
- Advisors: Diyi Yang, Tom Kwiatkowski

Education

Ph.D. Computing Science, University of Alberta   Edmonton, Canada
Thesis: Mitigating Hallucinations in Conversational LLMs.   2018 - 2022
- Advisor: Osmar Zaiane, GPA: 4.00/4.00

MSc. Computing Science, University of Alberta   Edmonton, Canada
Thesis: Response Generation For An Open-Ended Conversational Agent   2016 - 2018
- Advisor: Osmar Zaiane, GPA: 4.00/4.00

Preprints

You can find an exhaustive list of my publications in my Google Scholar profile.

1. **RL Grokking Recipe: How Does RL Unlock and Transfer New Algorithms in LLMs?.**
   Yiyou Sun, Yuhan Cao, Pohao Huang, Haoyue Bai, Hannaneh Hajishirzi, Dawn Song*,
   <u>Nouha Dziri</u>* (* = equal advising role)
   Arxiv 2025 (in submission to ICLR)

2. **Climbing the Ladder of Reasoning: What LLMs Can-and Still Can't-Solve after SFT?.**
   Yiyou Sun, Shawn Hu, Georgia Zhou, Ken Zheng, Hannaneh Hajishirzi, <u>Nouha Dziri</u>,
   Dawn Song.
   Arxiv 2025

3. **OpenAgentSafety: A Comprehensive Framework for Evaluating Real-World AI Agent
   Safety.** Sanidhya Vijayvargiya, Aditya Bharat Soni, Xuhui Zhou, Zora Zhiruo Wang,
   <u>Nouha Dziri</u>, Graham Neubig, Maarten Sap.
   Arxiv 2025 (in submission to ICLR)

4. **OMEGA: Can LLMs Reason Outside the Box in Math? Evaluating Exploratory, Compositional, and Transformative Generalization.** Yiyou Sun, Georgia Zhou, Hao Wang, Dacheng Li, Dawn Song*, <u>Nouha Dziri</u>* (* = equal advising role)
NeurIPS 2025

5. **Artificial Hivemind: The Open-Ended Homogeneity of Language Models (and Beyond)** Liwei Jiang, Yuanjun Chai, Margaret Li, Mickel Liu, Raymond Fok, Maarten Sap, Yulia Tsvetkov, <u>Nouha Dziri</u>, Yejin Choi.
NeurIPS 2025
Oral presentation (Top 1.4% among 25k submissions)

6. **Why and How LLMs Hallucinate: Connecting the Dots with Subsequence Associations** Yiyou Sun, Yu Gai, Lijie Chen, Abhi Ravichander, Yejin Choi, <u>Nouha Dziri</u>, Dawn Song.
NeurIPS 2025

7. **2 OLMo 2 Furious**. Pete Walsh, Luca Soldaini, Dirk Groeneveld, Kyle Lo, Shane Arora, ..., <u>Nouha Dziri</u>, Noah A. Smith, Hannaneh Hajishirzi.
COLM 2025
2M downloads on HuggingFace

8. **TÜLU 3: Pushing Frontiers in Open Language Model Post-Training**. Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze Brahman, Lester James V. Miranda, Alisa Liu, <u>Nouha Dziri</u>, Shane Lyu, Yuling Gu, Saumya Malik, Victoria Graf, Jena D. Hwang, Jiangjiang Yang, Ronan Le Bras, Oyvind Tafjord, Chris Wilhelm, Luca Soldaini, Noah A. Smith, Yizhong Wang, Pradeep Dasigi and Hannaneh Hajishirzi.
COLM 2025
300K downloads on HuggingFace

9. **SafetyAnalyst: Interpretable, Transparent, and Steerable LLM Safety Moderation** Jing-Jing Li, Valentina Pyatkin, Max Kleiman-Weiner, Liwei Jiang, <u>Nouha Dziri</u>, Anne GE Collins, Jana Schaich Borg, Maarten Sap, Yejin Choi, Sydney Levine.
ICML 2025

10. **AI as Humanity's Salieri: Quantifying Linguistic Creativity of Language Models via Systematic Attribution of Machine Text against Web Text**. Ximing Lu, Melanie Sclar, Skyler Hallinan, Niloofar Mireshghallah, Jiacheng Liu, Seungju Han, Allyson Ettinger, Liwei Jiang, Khyathi Chandu, <u>Nouha Dziri</u>, Yejin Choi.
ICLR 2025
Oral presentation (Top 1.8% among 12k submissions)
Featured in Science Magazine

11. **WildBench: Benchmarking LLMs with Challenging Tasks from Real Users in the Wild.** Bill Yuchen Lin, Yuntian Deng, Khyathi Chandu, Faeze Brahman, Abhilasha Ravichander, Valentina Pyatkin, <u>Nouha Dziri</u>, Ronan Le Bras, Yejin Choi.
ICLR 2025

12. **Steering Masked Discrete Diffusion Models Via Discrete Denoising Posterior Prediction** Jarrid Rector-Brooks, Mohsin Hasan, Zhangzhi Peng, Zachary Quinn, Chenghao Liu, Sarthak Mittal, <u>Nouha Dziri</u>, Michael Bronstein, Yoshua Bengio, Pranam Chatterjee, Alexander Tong, Avishek Joey Bose.
ICLR 2025

13. **RewardBench: Evaluating Reward Models for Language Modeling.** Nathan Lambert, Valentina Pyatkin, Jacob Morrison, LJ Miranda, Bill Yuchen Lin, Khyathi Chandu, <u>Nouha Dziri</u>, Sachin Kumar, Tom Zick, Yejin Choi, Noah A Smith, Hannaneh Hajishirzi.
NAACL 2025

14. **Rel-AI: An Interaction-Centered Approach To Measuring Human-LM Reliance.** Kaitlyn Zhou, Jena D Hwang, Xiang Ren, <u>Nouha Dziri</u>, Dan Jurafsky, Maarten Sap.
NAACL 2025
Best Paper Runner-Up 2025 (Top 0.2%)

15. **To Err is AI: A Case Study Informing LLM Flaw Reporting Practices** Sean McGregor, Allyson Ettinger, Nick Judd, Paul Albee, Liwei Jiang, Kavel Rao, Will Smith, Shayne Longpre, Avijit Ghosh, Christopher Fiorelli, Michelle Hoang, Sven Cattell, <u>Nouha Dziri</u>.
AAAI 2025

16. **WildTeaming at Scale: From In-the-Wild Jailbreaks to (Adversarially) Safer Language Models.** Liwei Jiang, Kavel Rao, Seungju Han, Allyson Ettinger, Faeze Brahman, Sachin Kumar, Niloofar Mireshghallah, Maarten Sap, Yejin Choi, <u>Nouha Dziri</u>.
NeurIPS 2024
Featured in MarkTechPost

17. **WildGuard: Open One-Stop Moderation Tools For Safety Risks, Jailbreaks, and Refusals of LLMs.** Seungju Han, Kavel Rao, Allyson Ettinger, Liwei Jiang, Bill Yuchen Lin, Nathan Lambert, Yejin Choi, and <u>Nouha Dziri</u>.
NeurIPS 2024
215K downloads on HuggingFace

18. **Multi-Attribute Constraint Satisfaction via Language Model Rewriting.** Ashutosh Baheti, Debanjana Chakraborty, Faeze Brahman, Ronan Le Bras, Ximing Lu, <u>Nouha Dziri</u>, Yejin Choi, Mark Riedl, Maarten Sap.
TMLR 2024

19. **A Roadmap to Pluralistic Alignment.** Taylor Sorensen, Jared Moore, Jillian Fisher, Mitchell Gordon, Niloofar Mireshghallah, Christopher Michael Rytting, Andre Ye, Liwei Jiang, Ximing Lu, <u>Nouha Dziri</u>, Tim Althoff, Yejin Choi.
ICML 2024

20. **Elastic Weight Removal For Faithful and Abstractive Dialogue generation.** Nico Daheim, <u>Nouha Dziri</u>, Mrinmaya Sachan, Iryna Gurevych, Edoardo Ponti.
NAACL 2024

21. **Faith and Fate: Limits of Transformers on Compositionality.**
<u>Nouha Dziri</u>, Ximing Lu, Melanie Sclar, Xiang Lorraine Li, Liwei Jian, Bill Yuchen Lin, Peter West, Chandra Bhagavatula, Ronan Le Bras, Jena D Hwang, Soumya Sanyal, Sean Welleck, Xiang Ren, Allyson Ettinger, Zaid Harchaoui, Yejin Choi.
NeurIPS 2023
Spotlight (Top 2% among 12k submissions)
Featured in Quanta Magazine, Science News, Montreal AI Ethics Institute
15 invited talks, 1 podcast, 2 guest lectures

22. **Fine-Grained Human Feedback Gives Better Rewards for Language Model Training.**
Zeqiu Wu, Yushi Hu, Weijia Shi, <u>Nouha Dziri</u>, Alane Suhr, Prithviraj Ammanabrolu, Noah A Smith, Mari Ostendorf, Hannaneh Hajishirzi.
NeurIPS 2023
Spotlight (Top 2% among 12k submissions)

23. **Self-Refine: Iterative Refinement with Self-Feedback.** Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegreffe, Uri Alon, <u>Nouha Dziri</u>, Shrimai Prabhumoye, Yiming Yang, Sean Welleck, Bodhisattwa Prasad Majumder, Shashank Gupta, Amir Yazdanbakhsh, Peter Clark.
NeurIPS 2023

24. **The Generative AI Paradox:" What It Can Create, It May Not Understand".** Peter West*, Ximing Lu*, <u>Nouha Dziri</u>*, Faeze Brahman*, Linjie Li*, Jena D Hwang, Liwei Jiang, Jillian Fisher, Abhilasha Ravichander, Khyathi Chandu, Benjamin Newman, Pang Wei Koh, Allyson Ettinger, Yejin Choi. (* = equal contribution)
ICLR 2024

25. **Phenomenal Yet Puzzling: Testing Inductive Reasoning Capabilities of Language Models with Hypothesis Refinement.**
Linlu Qiu, Liwei Jiang, Ximing Lu, Melanie Sclar, Valentina Pyatkin, Chandra Bhagavatula, Bailin Wang, Yoon Kim, Yejin Choi, <u>Nouha Dziri</u>*, Xiang Ren*. (* = equal advising role)
ICLR 2024
Oral presentation (Top 1.8% among 7k submissions)

26. **The Unlocking Spell on Base LLMs: Rethinking Alignment via In-Context Learning.**
Bill Yuchen Lin, Abhilasha Ravichander, Ximing Lu, <u>Nouha Dziri</u>, Melanie Sclar, Khyathi Chandu, Chandra Bhagavatula, Yejin Choi.
ICLR 2024

JOURNAL & CONFERENCE PUBLICATIONS

27. **Value Kaleidoscope: Engaging AI with pluralistic human values, rights, and duties**
Taylor Sorensen, Liwei Jiang, Jena Hwang, Sydney Levine, Valentina Pyatkin, Peter West, <u>**Nouha Dziri**</u>, Ximing Lu, Kavel Rao, Chandra Bhagavatula, Maarten Sap, John Tasioulas, Yejin Choi.
AAAI 2024

28. **Culture-Gen: Revealing global cultural perception in language models through natural language prompting**
Huihan Li, Liwei Jiang, Jena D Hwang, Hyunwoo Kim, Sebastin Santy, Taylor Sorensen, Bill Yuchen Lin, <u>**Nouha Dziri**</u>, Xiang Ren, Yejin Choi.
COLM 2024

29. **What Makes it Ok to Set a Fire? Iterative Self-distillation of Contexts and Rationales for Disambiguating Defeasible Social and Moral Situations** Kavel Rao, Liwei Jiang, Valentina Pyatkin, Yuling Gu, Niket Tandon, <u>**Nouha Dziri**</u>, Faeze Brahman, Yejin Choi.
EMNLP 2023

30. **Evaluating Open-Domain Question Answering in the Era of Large Language Models.**
Ehsan Kamalloo, <u>**Nouha Dziri**</u>, Charles LA Clarke, Davood Rafiei.
ACL 2023

31. **CHAMPAGNE: Learning Real-world Conversation from Large-Scale Web Videos.**
Seungju Han, Jack Hessel, <u>**Nouha Dziri**</u>, Yejin Choi, Youngjae Yu.
ICCV 2023

32. **On the Origin of Hallucinations in Conversational Models: Is it the Datasets or the Models?** <u>**Nouha Dziri**</u>, Sivan Milton, Mo Yu, Osmar Zaiane, Siva Reddy.
NAACL 2022

33. **FaithDial: A Faithful Benchmark for Information-Seeking Dialogue.**
<u>**Nouha Dziri**</u>, Ehsan Kamalloo, Sivan Milton, Osmar Zaiane, Mo Yu, Edoardo M Ponti, Siva Reddy.
TACL 2022

34. **Neural Path Hunter: Reducing Hallucination in Dialogue Systems via Path Grounding.**
<u>**Nouha Dziri**</u>, Andrea Madotto, Osmar Zaiane, Avishek Joey Bose.
EMNLP 2021

35. **Decomposed Mutual Information Estimation For Contrastive Representation Learning.**
Alessandro Sordoni*, <u>**Nouha Dziri**</u>*, Hannes Schulz*, Geoff Gordon, Philip Bachman, Remi Tachet Des Combes. (* = equal contribution)
ICML 2021

36. **Evaluating Coherence in Dialogue Systems Using entailment.**
<u>**Nouha Dziri**</u>*, Ehsan Kamalloo*, Kory W Mathewson, Osmar Zaiane. (* = equal contribution)
NAACL 2019

INVITED PANELS

Computer Use Agents
ICML workshop, Vancouver July 2025
Panelists: Graham Neubig, Ruslan Salakhutdinov, Yu Su, Victor Zhong, Alexandre Drouin

Data in Generative Models
ICML workshop, Vancouver July 2025
Panelists: Tatsunori Hashimoto, Eric Wong, Serena Booth, Pin-Yu Chen, Ivan Evtimov

AI Reasoning
Cross Future AI & Technology Summit, Vancouver July 2025
Panelists: Wenhu Chen, Ken Perlin

Trustworthy Foundation Models
International Symposium on Trustworthy Foundation Models, UAE May 2025
Panelists: Tomas Mikolov, Hakim Hacid, Tongliang Liu

INVITED PANELS	**Making the Right Career Call across Academia and Industry**	
	ICLR Workshop Women in ML, Singapore	April 2025
	Panelists: Claire Vernade, Reyhane Askari, Katherine Driscoll	
	Meta-Generation Algorithms for Large Language Models	
	NeurIPS Tutorial, Vancouver	Dec 2024
	Panelists: Noam Brown, Jakob Foerster, Rishabh Agarwal, Beidi Chen	
	Analyzing the Plasticity of LLMs During Language Interactions.	
	NeurIPS Language Gamification Workshop, Vancouver	Dec 2024
	Panelists: Aaron Courville, Alane Suhr, Tom Schaul, Marc Lanctot, Tom Griffiths, and Sam Devlin	
	Towards A Global Standard for AI Product Safety	
	MLCommons, San Francisco	Nov 2024
	Panelists: Ion Stoica, April Chen, Wan Sie Lee	

INVITED TALKS	**Reasoning in LLMs**	
	Workshop on Cognitive Basis of Reasoning/IVADO, Montreal	Jan 2026
	Autonomous LLM Agents: Risks and Scientific Challenges	
	IVADO/Mila - Quebec Artificial Intelligence Institute, Montreal	Nov 2025
	Faith and Fate: Limits of Transformers in Reasoning	
	D.E.Shaw Research, NYC	Sept 2025
	LLM Reasoning: Advanced Inference-time Strategies	
	Armenia LLM Summer School, Armenia	July 2025
	Can LLMs Reason Outside the Box in Math?	
	Apple Workshop on Reasoning and Planning, Cupertino	July 2025
	Faith and Fate: Limits of Transformers in Reasoning	
	Cross Future AI & Technology Summit, Vancouver	July 2025
	Trustworthy LLMs: How Data Quality Shapes Performance and Where It Falls Short?	
	ICML 2025 Workshop on Data in Generative Models, Vancouver	July 2025
	OpenAgentSafety: A Comprehensive Framework for Evaluating AI Agent Safety	
	ICML 2025 Workshop on Computer-Use Agents, Vancouver	July 2025
	Does Scaling Guarantee Trustworthy LLMs	
	International Symposium on Trustworthy Foundation Models, UAE	May 2025
	In-Context Learning in LLMs: Potential and Limits	
	Causality in the Era of Foundation Models Workshop, Barbados	Feb 2025
	Red-teaming and Safeguarding LLMs	
	International Association for Safe and Ethical AI, Paris	Feb 2025
	In-Context Learning in LLMs: Potential and Limits	
	NeurIPS Language Gamification Workshop, Vancouver	Dec 2024
	What it can create, it may not understand: Studying the Limits of Transformers	
	University of Cambridge	May 2024
	Limits of Generative AI Models and their Societal Implications.	
	Princeton University	Dec 2023
	Faith and Fate: Limits of Transformers on Compositionality	
	The Alan Turing Institute, UK	Nov 2023

INVITED TALKS	**Faith and Fate: Limits of Transformers on Compositionality.** *University of Edinburgh*	Nov 2023
	Faith and Fate: Limits of Transformers on Compositionality *SAIL workshop on fundamental limits of LLMs, Germany*	Oct 2023
	Faith and Fate: Limits of Transformers on Compositionality *University of Pittsburgh*	Oct 2023
	Faith and Fate: Limits of Transformers on Compositionality *Formal Languages and Neural Networks Seminar, US*	Sep 2023
	Towards Building Hallucination-Free Conversational Models *Stanford University*	Aug 2022
	FaithDial: A Faithful Benchmark for Information-Seeking Dialogue *Google Research, NYC*	May 2022
	FaithDial: A Faithful Benchmark for Information-Seeking Dialogue *Amazon Research, Seattle*	June 2022
	Evaluating Coherence in Dialogue Systems Using Entailment. *Google DeepMind, Montreal*	Dec 2019
GUEST LECTURER	**LLM Summer School 2025** ***Organizers:*** *ServiceNow Research/ Nvidia/ Meta* *Yerevan, Armenia* ***Lecture:*** *LLM Reasoning: Advanced Inference-Time Strategies*	*Summer 2025*
	11-430/830 Ethics, Safety, and Social Impact in NLP and LLMs ***Instructor:*** *Maarten Sap* *Carnegie Mellon University* ***Lecture:*** *Redteaming and Safegarding in LLMs*	*Winter 2025*
	Generative AI Seminar Course ***Instructor:*** *Adji Bousso Dieng* *Princeton University* ***Lecture:*** *Limits of Generative AI Models and their Societal Implications.*	*Fall 2023*
	CSE 599 D1: (Grad) Exploration on Language, Knowledge, and Reasoning ***Instructor:*** *Yejin Choi* *University of Washington* ***Lecture:*** *Can LLMs reason?*	*Winter 2023*
	CMPUT 650: (Grad) Computational Semantics ***Instructor:*** *Greg Kondrak* *University of Alberta* ***Lecture:*** *Challenges in Conversational LLMs*	*Winter 2018*
TEACHING ASSISTANT	**CMPUT 101 - Introduction to Computing** *University of Alberta* Instructor: Marianne Morris	2018

RESEARCH BLOG POST

- *Have o1 Models Cracked Human Reasoning?*, 2024.
- *Current Paradigms of LLMs Safety Alignment are superficial*, 2024.
- *DeepSeek R1: Innovative Research and Engineering Can Rival Brute-Force Scaling*, 2025.
- *Can LLMs Reason Outside the Box in Math?*, 2025.
- *RL Grokking Recipe – How Can We Enable LLMs to Solve Previously Unsolvable Tasks with RL?*, 2025.

STUDENTS MENTORING

- **Ximing Lu**, Predoctoral student at Ai2 –> PhD student at UW   2022-2024
  ICLR 2025 Oral (top 1%), NeurIPS 2023 Spotlight (top 2%)
- **Liwei Jiang**, PhD student at UW   2023-2025
  NeurIPS 2025 Oral (top 1%), NeurIPS 2024
- **Seungju Han**, BSc at Seoul National University -> PhD at Stanford   2023-2024
  NeurIPS 2024, ICCV 2023
- **Kavel Rao**, BSc student at UW -> SWE at Jane Street   2023-2024
  NeurIPS 2024, EMNLP 2023
  Single Awardee of the 2024 Best Senior Thesis Award at UW CSE
- **Linlu Qiu**, PhD student at MIT   2023
  ICLR 2024 Oral (top 1%)
- **Huihan Li**, PhD student at University of Southern California   2023-2024
  COLM 2024
- **Yiyou Sun**, Postdoc at University of California, Berkeley   2025-current
  2 papers at NeurIPS 2025
- **Shrey Jain**, MSc student at Carnegie Mellon University   2025-current

WORKSHOP CO-ORGANIZER

- The first workshop for Research on Agent Language Models (REALM)   ACL 2025
- The Reasoning and Planning for Large Language Models   ICLR 2025
- System 2 Reasoning at Scale   NeurIPS 2024
- The third workshop on Document-Grounded Dialogue Systems (DialDoc)   ACL 2023

ACADEMIC SERVICES

Demo Chair: NAACL 2025

Senior Area Chair: ACL 2025 in the area of Ethics, Bias, and Fairness

Area Chair: EMNLP 2023, COLM (2024-2025)

Reviewer: NeurIPS (2022-2024), ICLR (2022-2024), ACL (2018-2023), EMNLP (2018-2023), NAACL (2018-2023), EACL (2018-2022)

PRESS

- **Quanta Magazine**: *Chatbot Software Begins to Face Fundamental Limitations*, 2025.
- **Science Magazine**: *AI writing is improving, but it can't match human creativity*, 2024.
- **LeMonde**: *Should we be concerned about the 'hallucinations' of AI like ChatGPT or Gemini?*, 2024.
- **ScienceNews**: *AI's reasoning skills can't be assessed by current tests*, 2024.
- **TechCrunch**: *Treating a chatbot nicely might boost its performance — here's why*, 2024.
- **Podcast (Women in AI Research)**: *Limits of Transformers*, 2025.
- **MarkTechPost**: *An Automatic Red-Team Framework to Compose Adversarial Attacks*, 2024.

AWARDS
AND
HONORS

- **Runner-up Best Paper** at NAACL 2025 **(top 0.1%)**   2025
- **Outstanding Reviewer** at ACL 2021 **(top 1%)**   2021
- **Alberta Doctoral Recruitment Scholarship** ($10,000)   2018
- **Mitacs Globalink PhD Graduate Fellowship** ($44,000)   2018
- **National Scholarship for MSc and PhD studies** ($70,000)   2016
- **Best Poster Award,** ACM Canadian Celebration of Women in Computing ($400)   2017
- **Mitacs Globalink MSc Graduate Fellowship** ($15,000)   2016
- **DAAD Scholarship** for Research Internship, Leipzig, Germany (€10.000)   2015
- **Erasmus Mundus Exchange Scholarship** for BSc studies(€50.000)   2015